

In [65]:

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
import sklearn
import matplotlib.pyplot as plt
import seaborn as sns
import hvplot.pandas
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LinearRegression
##cross validation
from sklearn.model_selection import cross_val_score
```

In [28]:

```
heart_data = pd.read_csv('C:/Users/vikas pawar/Downloads/heart_disease_data.csv')
```

In [29]:

```
heart_data.head()
```

Out[29]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	targ
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	

In [30]:

```
heart_data.tail()
```

Out[30]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	ta
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	

In [31]:

```
heart_data.shape
```

Out[31]:

(303, 14)

In [32]:

```
heart_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trestbps    303 non-null   int64
4   chol        303 non-null   int64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalach     303 non-null   int64
8   exang       303 non-null   int64
9   oldpeak     303 non-null   float64
10  slope       303 non-null   int64
11  ca          303 non-null   int64
12  thal        303 non-null   int64
13  target      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

In [33]:

```
heart_data.isnull().sum()
```

Out[33]:

```
age          0
sex          0
cp           0
trestbps     0
chol         0
fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64
```

In [34]:

```
heart_data.describe()
```

Out[34]:

	age	sex	cp	trestbps	chol	fbs	restecg
<b>count</b>	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
<b>mean</b>	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053
<b>std</b>	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860
<b>min</b>	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000
<b>25%</b>	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000
<b>50%</b>	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000
<b>75%</b>	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000
<b>max</b>	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000

In [35]:

```
heart_data['target'].value_counts()
```

Out[35]:

```
1    165
0    138
Name: target, dtype: int64
```

In [36]:

```
X = heart_data.drop(columns='target', axis=1)
Y = heart_data['target']
```

In [37]:

```
print(X)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
\										
0	63	1	3	145	233	1	0	150	0	2.3
1	37	1	2	130	250	0	1	187	0	3.5
2	41	0	1	130	204	0	0	172	0	1.4
3	56	1	1	120	236	0	1	178	0	0.8
4	57	0	0	120	354	0	1	163	1	0.6
..	...	...	..	...	...	...	...	...	...	...
298	57	0	0	140	241	0	1	123	1	0.2
299	45	1	3	110	264	0	1	132	0	1.2
300	68	1	0	144	193	1	1	141	0	3.4
301	57	1	0	130	131	0	1	115	1	1.2
302	57	0	1	130	236	0	0	174	0	0.0

	slope	ca	thal
0	0	0	1
1	0	0	2
2	2	0	2
3	2	0	2
4	2	0	2
..	...	..	...
298	1	0	3
299	1	0	3
300	1	2	3
301	1	1	3
302	1	1	2

[303 rows x 13 columns]

In [38]:

```
print(Y)
```

0	1
1	1
2	1
3	1
4	1
..	
298	0
299	0
300	0
301	0
302	0

Name: target, Length: 303, dtype: int64

In [39]:

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=1)
```

In [40]:

```
print(X.shape, X_train.shape, X_test.shape)
```

(303, 13) (242, 13) (61, 13)

In [41]:

```
model = LogisticRegression()
```

In [42]:

```
model.fit(X_train, Y_train)
```

C:\Users\vikas pawar\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

Out[42]:

```
LogisticRegression()
```

In [43]:

```
X_train_prediction = model.predict(X_train)  
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

In [44]:

```
print('Accuracy on Training data : ', training_data_accuracy)
```

Accuracy on Training data : 0.8512396694214877

In [45]:

```
X_test_prediction = model.predict(X_test)  
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

In [46]:

```
print('Accuracy on Test data : ', test_data_accuracy)
```

Accuracy on Test data : 0.819672131147541

In [47]:

```

input_data = (62,0,0,140,268,0,0,160,0,3.6,0,2,2)

# change the input data to a numpy array
input_data_as_numpy_array= np.asarray(input_data)

# reshape the numpy array as we are predicting for only on instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = model.predict(input_data_reshaped)
print(prediction)

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')

```

[0]

The Person does not have a Heart Disease

C:\Users\vikas pawar\anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

```
warnings.warn(
```

In [49]:

heart\_data.mean

Out[49]:

```

<bound method NDFrame._add_numeric_operations.<locals>.mean of
ex cp trestbps chol fbs restecg thalach exang oldpeak \
0 63 1 3 145 233 1 0 150 0 2.3
1 37 1 2 130 250 0 1 187 0 3.5
2 41 0 1 130 204 0 0 172 0 1.4
3 56 1 1 120 236 0 1 178 0 0.8
4 57 0 0 120 354 0 1 163 1 0.6
.. ... ..
298 57 0 0 140 241 0 1 123 1 0.2
299 45 1 3 110 264 0 1 132 0 1.2
300 68 1 0 144 193 1 1 141 0 3.4
301 57 1 0 130 131 0 1 115 1 1.2
302 57 0 1 130 236 0 0 174 0 0.0

```

```

slope ca thal target
0 0 0 1 1
1 0 0 2 1
2 2 0 2 1
3 2 0 2 1
4 2 0 2 1
.. ... ..
298 1 0 3 0
299 1 0 3 0
300 1 2 3 0
301 1 1 3 0
302 1 1 2 0

```

[303 rows x 14 columns]&gt;

In [50]:

```
heart_data.mode
```

Out[50]:

<bound method DataFrame.mode of						age	sex	cp	trestbps	chol	fbs	re
stecg thalach exang oldpeak \												
0	63	1	3	145	233	1		0	150	0		2.3
1	37	1	2	130	250	0		1	187	0		3.5
2	41	0	1	130	204	0		0	172	0		1.4
3	56	1	1	120	236	0		1	178	0		0.8
4	57	0	0	120	354	0		1	163	1		0.6
..	...	...	..	...	...	...		...	...	...		...
298	57	0	0	140	241	0		1	123	1		0.2
299	45	1	3	110	264	0		1	132	0		1.2
300	68	1	0	144	193	1		1	141	0		3.4
301	57	1	0	130	131	0		1	115	1		1.2
302	57	0	1	130	236	0		0	174	0		0.0

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..	...	..	...	...
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]>

In [51]:

```
heart_data.notnull()
```

Out[51]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	ti
0	True	True	True	True	True	True	True	True	True	True	True	True	True
1	True	True	True	True	True	True	True	True	True	True	True	True	True
2	True	True	True	True	True	True	True	True	True	True	True	True	True
3	True	True	True	True	True	True	True	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True	True	True	True	True	True
...	...	...	...	...	...	...	...	...	...	...	...	...	...
298	True	True	True	True	True	True	True	True	True	True	True	True	True
299	True	True	True	True	True	True	True	True	True	True	True	True	True
300	True	True	True	True	True	True	True	True	True	True	True	True	True
301	True	True	True	True	True	True	True	True	True	True	True	True	True
302	True	True	True	True	True	True	True	True	True	True	True	True	True

303 rows × 14 columns

In [52]:

```
heart_data.values
```

Out[52]:

```
array([[63., 1., 3., ..., 0., 1., 1.],
       [37., 1., 2., ..., 0., 2., 1.],
       [41., 0., 1., ..., 0., 2., 1.],
       ...,
       [68., 1., 0., ..., 2., 3., 0.],
       [57., 1., 0., ..., 1., 3., 0.],
       [57., 0., 1., ..., 1., 2., 0.]])
```



In [53]:

```
heart_data.isin([0]).any()  
(heart_data==0).sum()
```

Out[53]:


```
age          0  
sex          96  
cp          143  
trestbps     0  
chol         0  
fbs         258  
restecg     147  
thalach      0  
exang       204  
oldpeak     99  
slope       21  
ca         175  
thal        2  
target     138  
dtype: int64
```

In [59]:

```

import seaborn as sns
#get correlations of each features in dataset
corrmat = heart_data.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(20,20))
#plot heat map
g=sns.heatmap(heart_data[top_corr_features].corr(),annot=True,cmap="RdYlGn")

```




In [60]:

```
heart_data.hist()
```

Out[60]:

```
array([[<AxesSubplot:title={'center':'age'}>,
       <AxesSubplot:title={'center':'sex'}>,
       <AxesSubplot:title={'center':'cp'}>,
       <AxesSubplot:title={'center':'trestbps'}>],
       [<AxesSubplot:title={'center':'chol'}>,
       <AxesSubplot:title={'center':'fbs'}>,
       <AxesSubplot:title={'center':'restecg'}>,
       <AxesSubplot:title={'center':'thalach'}>],
       [<AxesSubplot:title={'center':'exang'}>,
       <AxesSubplot:title={'center':'oldpeak'}>,
       <AxesSubplot:title={'center':'slope'}>,
       <AxesSubplot:title={'center':'ca'}>],
       [<AxesSubplot:title={'center':'thal'}>,
       <AxesSubplot:title={'center':'target'}>, <AxesSubplot:>,
       <AxesSubplot:>]], dtype=object)
```



In [61]:

```
y = heart_data['target']
X = heart_data.drop(['target'], axis = 1)
```

In [62]:

```
from sklearn.model_selection import cross_val_score
knn_scores = []
for k in range(1,21):
    knn_classifier = KNeighborsClassifier(n_neighbors = k)
    score=cross_val_score(knn_classifier,X,y,cv=10)
    knn_scores.append(score.mean())
```

In [63]:

```
plt.plot([k for k in range(1, 21)], knn_scores, color = 'red')
for i in range(1,21):
    plt.text(i, knn_scores[i-1], (i, knn_scores[i-1]))
plt.xticks([i for i in range(1, 21)])
plt.xlabel('Number of Neighbors (K)')
plt.ylabel('Scores')
plt.title('K Neighbors Classifier scores for different K values')
```

Out[63]:

Text(0.5, 1.0, 'K Neighbors Classifier scores for different K values')

